

Postiz Integration for StorageAds

How we bundle social + Google Business Profile posting into the platform — the economics, the architecture, and the stack to make it run itself.

TL;DR

We bundle a 'social + GBP on autopilot' feature into StorageAds. Postiz (open-source) handles the publishing rail. We build the storage-specific brain on top: content engine, asset generation, response automation, and cross-portfolio learning. Clients never see Postiz. Agency workload trends toward zero. COGS lands around **\$3-5 per facility per month** with **90%+ gross margin** at typical bundling prices.

The honest framing: Postiz is roughly 15-20% of the system. The other 80% is what we build, and that's where the moat lives.

What Postiz Is (and Isn't)

Postiz is an open-source social media scheduler. It connects to Google Business Profile, Facebook, Instagram, LinkedIn, X, TikTok, Threads, and more via OAuth. It can be self-hosted (AGPL license, basically free) or run on their cloud.

What Postiz handles for us:

- OAuth into every social platform (the boring, expensive part)
- Scheduling and cross-posting across multiple accounts
- Multi-workspace organization (one workspace per facility)
- Basic publish-level analytics

What Postiz does NOT handle:

- Writing the posts
- Generating images or video
- Triaging comments or DMs
- Responding to reviews
- Knowing anything about storage facilities

That last gap is the opportunity. Generic social tools cannot match a storage-vertical-specific layer powered by our occupancy data, audit findings, and cross-portfolio learning.

The Vision — 6 Layers

Picture the whole thing as a stack. Data flows up; nothing in any layer above is in Postiz. Postiz only owns Layer 4.

Layer 1 — Signal (our unfair advantage)

Every facility we onboard already gives StorageAds more data than any standalone tool will ever have:

- Occupancy by unit size from FMS uploads
- Active promos, current pricing
- Competitor pricing + occupancy from our scraping
- Audit-tool findings (every audit = a content backlog)
- Reviews stream (Google + FB)
- Local context: weather, school move-ins, military PCS cycles, hurricane windows
- Past post performance across the whole portfolio

Layer 2 — Content Engine (the brain)

Not 'AI writes social posts.' A storage-specific orchestrator that decides what to post, when, and why. Examples of real decisions it makes:

- *Facility A has 3 climate 10x10s left and is 92% occupied → push an availability post today*
- *Facility B got a 5-star review yesterday → schedule a review-as-quote card for Saturday morning*
- *Hurricane forecast hitting FL panhandle in 6 days → push prep-storage content to every coastal facility*
- *Facility C's GBP hasn't posted in 9 days → engine auto-queues from the template library*

Calendar programs itself. Always 14 days scheduled ahead. Never an empty queue.

Layer 3 — Asset Layer (Angelo's domain)

Visuals can't look like generic AI. Clients need photos that look like *their* facility.

- Pull Google Maps + Street View + client uploads as base assets
- Brand them once per facility (logo, color, name overlay)
- Template-driven cards for availability / promo / review (dynamic data fills in)
- Runway pipeline does short B-roll reels from photo sets
- Auto-crop per platform (1:1 IG, 9:16 reel, 16:9 GBP, 4:5 FB)

One photo upload at onboarding generates 60 days of asset variants.

Layer 4 — Publish (Postiz earns its keep here)

- One workspace per facility under our master account
- OAuth all client accounts during the onboarding wizard
- Time-of-day optimization per platform per local timezone
- Approval mode configurable per client: full auto (default), agency-approved, or client-approved
- Bulk pause/push across the whole portfolio in one click (e.g. kill everything in a region during a crisis, push hurricane content everywhere on the coast)
- Direct GBP API fallback for the things Postiz doesn't cover well

Layer 5 — Response (the real margin saver)

Publishing is easy. Responding is what kills agency margin. This is the layer most social tools don't do and most agencies do badly.

- **GBP Q&A:** AI drafts the answer using facility data, we one-click approve. Or seed our own questions and answers.
- **Review responses:** AI drafts, we approve. We already have the Claude API plumbing for this.
- **FB/IG comments:** AI triages into question / spam / complaint / positive. Auto-responds to easy. Queues hard ones for human review.
- **DM auto-reply with intent detection:** 'do you have a 10x10' → push to call-tracking number and open a lead. 'your gate is broken' → escalate to facility manager. Spam → ignore.
- **Unified inbox in the StorageAds portal:** we work one queue instead of N platforms times M facilities.

Layer 6 — Learning (the moat)

Every post across every facility on our platform is training data Postiz alone can never use:

- Which content types drive calls vs. just likes? Reweight the engine.
- Best post times per metro, facility size, season
- Cross-portfolio benchmarks: *'your GBP is in the 73rd percentile of storage facilities in our network'*
- Audit tool tie-in: every audit run becomes the first 30 days of content for new clients
- Reporting that ties social to leases, not impressions: *'GBP drove 4 calls, 2 reservations, \$189 ARPU last month'*

Economics — What This Actually Costs Us

Real numbers per facility per month, fully loaded, assuming self-hosted Postiz and our existing AI pipeline:

Line item	Cost / facility / mo	Notes
Postiz infrastructure	\$0.10 - \$0.50	Hetzner box amortized across 100-500 workspaces
Anthropic API — post generation	\$0.20 - \$0.50	8-20 posts/mo on Sonnet 4.6
Anthropic API — responses + triage	\$0.50 - \$1.50	GBP Q&A, review replies, DM intent
Image generation (FAL)	\$0.20 - \$1.00	Inference cost only — pipeline already built
Video generation (Runway)	\$0.80 - \$2.00	Better/Best tiers only, ~2 reels/mo
GBP / Meta Graph APIs	\$0	Free
Database / Vercel functions	~\$0.10	Rounding error
TOTAL	\$2 - \$5	Full feature set, self-hosted

Margin math

Bake the feature into Better/Best as a \$39-79/mo add-on or use it to raise tier floors: **90%+ gross margin**. Even bundled into Good as a loss leader to make the tier impossible to comparison-shop, COGS stays under 5% of revenue.

Hard rule on hosting

Stay on Postiz Cloud for the pilot (fast, kill switch is easy). **Self-host the moment we go past ~10 paying clients**. Per-workspace cloud pricing eats margin alive at scale.

Scalability — When Things Actually Break

Real bottlenecks, in order of when they hit:

Client count	Bottleneck	What we do
~50	None	Single box, single dev. Cruise.
~250	OAuth re-auth fatigue (Meta tokens expire ~60 days)	Self-service re-auth flow in client portal. Build this early.
~500	Postiz queue workers	Horizontal scaling on Hetzner/Railway. Trivial.
~1000	Cross-portfolio analytics queries slow	Move analytics to DuckDB / ClickHouse / BigQuery. Not a 2026 problem.
~2000	Anthropic API tier limits	Tier 4+ should hold. Fall back to Haiku for cheap operations like triage.

Yes, we offer this to every client in every tier. Differentiate by post volume, video inclusion, and how autonomous the response layer is — not by whether the feature exists.

Recommended Stack Beyond Postiz

The tools that actually move the needle on quality and reducing our admin workload. Grouped by purpose.

Content + Intelligence

- **Anthropic Claude** (already wired) — orchestrator and writer
- **Tavily** or **Exa** — local news/event scanning for content triggers (~\$10-50/mo)
- **Perplexity API** — competitor monitoring per metro (~\$20/mo)
- **OpenWeather** or **Tomorrow.io** — weather-triggered content (free tier fine)
- **Whisper** or **Deepgram** — owner voice memos become posts. *'We just finished the new climate building'* becomes 5 scheduled posts.
- **Brandfetch API** — auto-pull client logo/colors at OAuth time (~\$30/mo)

Asset Generation

- **FAL.ai** (your domain) — image generation
- **Runway** (your domain) — video generation
- **Bannerbear** or **Placid** — templated card generation via API (~\$49-99/mo). Way faster than building our own template engine for availability / promo / review cards.

Workflow Engine (this cuts admin work the most)

- **Inngest** or **Trigger.dev** — durable workflows for the content pipeline. Better than raw Vercel cron for retry / backoff / observability. Free tier covers us to ~50 clients.
- **n8n** (self-hosted) — visual workflow builder for the long tail of 'when X happens, do Y' rules we'll accumulate. Self-host on the same box as Postiz.

Response Layer

- **Sendbird** or **Pusher** — real-time inbox without building WebSocket plumbing
- **Build the unified inbox native.** Don't outsource to Front or Missive — we lose the vertical context that makes the storage angle work.

Observability (so we keep our sanity)

- **Langfuse** or **Helicone** — LLM observability. Track cost / latency / quality per facility (\$0-99/mo).
- **Sentry** — error tracking (already in the task plan)
- **PostHog** — product analytics + session replay on the client portal so we can see why clients churn.

Client Portal Experience (reduces their workload too)

- **Loom Embed API** — weekly performance recap as a 90-second video instead of a report nobody reads
- **Cal.com** (have) — booking
- **Resend** (have) — transactional email

- **Knock** — notification routing across email / SMS / in-app / push so clients get pinged the right way for the right thing

The Compounding Play

If we wire all this up, every audit we run on a prospect produces 30 days of social content automatically the moment they sign. That's a sales weapon, not just a feature.

Sign tonight and your social calendar is full by Monday.

That line is unbeatable against any other agency pitch in storage.

Real numbers to deliver it:

- **\$3-5 COGS** per facility per month
- **Under 5 minutes** of agency labor per facility per week (exceptions only)
- **One ops person runs 500+ facilities** at the workload of 50 today

What to Flag Before We Commit

- **Content safety.** Full auto = one bad post away from a PR problem. We need a rules layer (no politics, no naming competitors, no pricing claims we can't honor, no AI-generated fake storage units).
- **GBP rate limits + post rejections** are real and not always predictable. Queue + retry + agency alert is mandatory.
- **Voice drift.** If every facility's content comes from the same engine, they start sounding alike. Each facility needs a per-facility voice profile (formal vs. casual, urban vs. rural, family-owned vs. corporate).
- **Postiz dependency.** If their roadmap diverges from ours (e.g. they deprioritize GBP), we're stuck. Thin adapter pattern means we can swap the publish layer in 2 days, not rewrite.

What I Need From You (Angelo)

1. Reality check on the asset layer — can your FAL + Runway pipelines produce per-facility variants at scale without us hand-prompting each one?
2. Opinion on hosting Postiz alongside our existing infra vs. a dedicated box
3. Anything in the response/inbox layer that conflicts with the ad platform work you're already doing
4. Veto power on any of this if you see a blocker I'm missing

Next Decisions to Make

- Pilot Postiz Cloud against one of my portfolio facilities to see end-to-end flow
- If it works, scope the self-hosted deployment + the thin StorageAds adapter
- Identify which 2-3 storage-specific content types we build templates for first
- Decide bundling strategy: Better/Best baked-in vs. dedicated add-on

Open to your read on this. Reply with reactions, objections, or pieces you'd want to own.